

```
import pandas as pd
import numpy as np

# till now
# Timestamp
pd.Timestamp('6th jan 2023 8:10')
# DatetimeIndex -> df and series index
pd.DatetimeIndex([pd.Timestamp('6th jan 2023 8:10'),pd.Timestamp('7th jan 2023 8:10'),pd.Timestamp('8th jan 2023 8:10')])[0]
# date_range()
pd.date_range(start='2023-1-6',end='2023-1-31',freq='D')
# to_datetime()
s = pd.Series(['2023/1/6','2023/1/7','2023/1/7'])
pd.to_datetime(s).dt.day_name()

0      Friday
1     Saturday
2     Saturday
dtype: object
```

## ▼ Timedelta Object

Represents a duration, the difference between two dates or times.

```
# create using Timestamp objects
t1 = pd.Timestamp('6th Jan 2023 8:20:14')
t2 = pd.Timestamp('26th Jan 2023 10:00:00')

t2 - t1
```

```
Timedelta('20 days 01:39:46')
```

```
# standalone creation
pd.Timedelta(days=2,hours=10,minutes=35)
```

```
Timedelta('2 days 10:35:00')
```

```
# Arithmetic
pd.Timestamp('6th jan 2023') + pd.Timedelta(days=2,hours=10,minutes=35)
```

```
Timestamp('2023-01-08 10:35:00')
```

```
pd.date_range(start='2023-1-6',end='2023-1-31',freq='D') - pd.Timedelta(days=2,hours=10,minutes=35)
```

```
DatetimeIndex(['2023-01-03 13:25:00', '2023-01-04 13:25:00',
                '2023-01-05 13:25:00', '2023-01-06 13:25:00',
                '2023-01-07 13:25:00', '2023-01-08 13:25:00',
                '2023-01-09 13:25:00', '2023-01-10 13:25:00',
                '2023-01-11 13:25:00', '2023-01-12 13:25:00',
                '2023-01-13 13:25:00', '2023-01-14 13:25:00',
                '2023-01-15 13:25:00', '2023-01-16 13:25:00',
                '2023-01-17 13:25:00', '2023-01-18 13:25:00',
                '2023-01-19 13:25:00', '2023-01-20 13:25:00',
                '2023-01-21 13:25:00', '2023-01-22 13:25:00',
                '2023-01-23 13:25:00', '2023-01-24 13:25:00',
                '2023-01-25 13:25:00', '2023-01-26 13:25:00',
                '2023-01-27 13:25:00', '2023-01-28 13:25:00'],
              dtype='datetime64[ns]', freq='D')
```

```
# real life example
df = pd.read_csv('deliveries.csv')
df.head()
```

|   | order_date | delivery_date |
|---|------------|---------------|
| 0 | 5/24/98    | 2/5/99        |
| 1 | 4/22/92    | 3/6/98        |
| 2 | 2/10/91    | 8/26/92       |
| 3 | 7/21/92    | 11/20/97      |
| 4 | 9/2/93     | 6/10/98       |

```
df['order_date'] = pd.to_datetime(df['order_date'])
df['delivery_date'] = pd.to_datetime(df['delivery_date'])
```

```
df['delivery_time_period'] = df['delivery_date'] - df['order_date']

df['delivery_time_period'].mean()

Timedelta('1217 days 22:53:53.532934128')
```

▼ Time series

A time series is a data set that tracks a sample over time. In particular, a time series allows one to see what factors influence certain variables from period to period. Time series analysis can be useful to see how a given asset, security, or economic variable changes over time.

Examples

- Financial Data (Company stocks)
- Natural Data (Rainfall measurement)
- Event Data (Covid)
- Medical Data (Heart rate monitoring)

Types of Operations done on Time Series

- Time Series Analysis
- Time Series Forecasting

```
google = pd.read_csv('google.csv')
google.head()
```

|   | Date       | Open      | High      | Low       | Close     | Adj Close | Volume   |
|---|------------|-----------|-----------|-----------|-----------|-----------|----------|
| 0 | 2004-08-19 | 49.813290 | 51.835709 | 47.800831 | 49.982655 | 49.982655 | 44871361 |
| 1 | 2004-08-20 | 50.316402 | 54.336334 | 50.062355 | 53.952770 | 53.952770 | 22942874 |
| 2 | 2004-08-23 | 55.168217 | 56.528118 | 54.321388 | 54.495735 | 54.495735 | 18342897 |
| 3 | 2004-08-24 | 55.412300 | 55.591629 | 51.591621 | 52.239197 | 52.239197 | 15319808 |
| 4 | 2004-08-25 | 52.284027 | 53.798351 | 51.746044 | 52.802086 | 52.802086 | 9232276  |

```
google.tail()
```

|      | Date       | Open        | High        | Low         | Close       | Adj Close   | Volume  |
|------|------------|-------------|-------------|-------------|-------------|-------------|---------|
| 4466 | 2022-05-16 | 2307.679932 | 2332.149902 | 2286.699951 | 2295.850098 | 2295.850098 | 1164100 |
| 4467 | 2022-05-17 | 2344.550049 | 2344.550049 | 2306.750000 | 2334.030029 | 2334.030029 | 1078800 |
| 4468 | 2022-05-18 | 2304.750000 | 2313.913086 | 2242.840088 | 2248.020020 | 2248.020020 | 1399100 |
| 4469 | 2022-05-19 | 2236.820068 | 2271.750000 | 2209.360107 | 2214.909912 | 2214.909912 | 1459600 |
| 4470 | 2022-05-20 | 2241.709961 | 2251.000000 | 2127.459961 | 2186.260010 | 2186.260010 | 1878100 |

```
subs = pd.read_csv('subscribers.csv')
views = pd.read_csv('views.csv')
wt = pd.read_csv('watch-time.csv')
comments = pd.read_csv('comments.csv')
```

```
comments.tail()
```

|      | Date       | Comments added |
|------|------------|----------------|
| 1354 | 2022-12-31 | 20             |
| 1355 | 2023-01-01 | 13             |
| 1356 | 2023-01-02 | 66             |
| 1357 | 2023-01-03 | 22             |
| 1358 | 2023-01-04 | 21             |

```
yt = subs.merge(views,on='Date').merge(wt,on='Date').merge(comments,on='Date')
```

```
yt.head()
```

|   | Date       | Subscribers | Views | Watch time (hours) | Comments added |
|---|------------|-------------|-------|--------------------|----------------|
| 0 | 2019-04-17 | 0           | 0     | 0.0                | 0              |
| 1 | 2019-04-18 | 0           | 0     | 0.0                | 0              |
| 2 | 2019-04-19 | 0           | 0     | 0.0                | 0              |
| 3 | 2019-04-20 | 0           | 0     | 0.0                | 0              |
| 4 | 2019-04-21 | 0           | 0     | 0.0                | 0              |

```
google['Date'] = pd.to_datetime(google['Date'])
yt['Date'] = pd.to_datetime(yt['Date'])
```

```
yt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1359 entries, 0 to 1358
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  1359 non-null  datetime64[ns]
1   Subscribers           1359 non-null  int64
2   Views                 1359 non-null  int64
3   Watch time (hours)    1359 non-null  float64
4   Comments added        1359 non-null  int64
dtypes: datetime64[ns](1), float64(1), int64(3)
memory usage: 63.7 KB
```

```
google.set_index('Date',inplace=True)
yt.set_index('Date',inplace=True)
```

```
yt.head()
```

|            | Subscribers | Views | Watch time (hours) | Comments added |
|------------|-------------|-------|--------------------|----------------|
| Date       |             |       |                    |                |
| 2019-04-17 | 0           | 0     | 0.0                | 0              |
| 2019-04-18 | 0           | 0     | 0.0                | 0              |
| 2019-04-19 | 0           | 0     | 0.0                | 0              |
| 2019-04-20 | 0           | 0     | 0.0                | 0              |
| 2019-04-21 | 0           | 0     | 0.0                | 0              |

```
# fetch a specific date
google.loc['2021-12-30']
```

```
Open          2929.000000
High          2941.250000
Low           2915.169922
Close         2920.050049
Adj Close     2920.050049
Volume        648900.000000
Name: 2021-12-30 00:00:00, dtype: float64
```

```
# partial indexing -> select a particular year/month
yt.loc['2022-12']
yt.loc['2022']
```

|            | Subscribers | Views | Watch time (hours) | Comments added |
|------------|-------------|-------|--------------------|----------------|
| Date       |             |       |                    |                |
| 2022-01-01 | 30          | 3111  | 327.3357           | 11             |
| 2022-01-02 | 42          | 4109  | 402.5877           | 5              |
| 2022-01-03 | 45          | 4264  | 431.8496           | 13             |
| 2022-01-04 | 53          | 4160  | 427.6701           | 7              |
| 2022-01-05 | 32          | 3971  | 433.6201           | 8              |
| ...        | ...         | ...   | ...                | ...            |
| 2022-12-27 | 201         | 18245 | 2228.4891          | 18             |

```
# create year month-name and day-name
yt['month_name'] = yt.index.month_name()
yt['weekday_name'] = yt.index.day_name()
yt['quarter'] = yt.index.quarter
```

yt.head()

|            | Subscribers | Views | Watch time (hours) | Comments added | month_name | weekdday_name | quarter | weekday_name |
|------------|-------------|-------|--------------------|----------------|------------|---------------|---------|--------------|
| Date       |             |       |                    |                |            |               |         |              |
| 2019-04-17 | 0           | 0     | 0.0                | 0              | April      | April         | 2       | Wednesday    |
| 2019-04-18 | 0           | 0     | 0.0                | 0              | April      | April         | 2       | Thursday     |
| 2019-04-19 | 0           | 0     | 0.0                | 0              | April      | April         | 2       | Friday       |

```
google['month_name'] = google.index.month_name()
google['weekday_name'] = google.index.day_name()
google['quarter'] = google.index.quarter

google.head()
```

|            | Open      | High      | Low       | Close     | Adj Close | Volume   | month_name | weekday_name | quarter |
|------------|-----------|-----------|-----------|-----------|-----------|----------|------------|--------------|---------|
| Date       |           |           |           |           |           |          |            |              |         |
| 2004-08-19 | 49.813290 | 51.835709 | 47.800831 | 49.982655 | 49.982655 | 44871361 | August     | Thursday     | 3       |
| 2004-08-20 | 50.316402 | 54.336334 | 50.062355 | 53.952770 | 53.952770 | 22942874 | August     | Friday       | 3       |
| 2004-08-23 | 55.168217 | 56.528118 | 54.321388 | 54.495735 | 54.495735 | 18342897 | August     | Monday       | 3       |

```
# slicing
yt.loc['2022-12-15': '2023-1-1':2]
```


```
# challenge -> fetch info for a particular date every year -> limitation of timedelta
google.head()

google[google.index.isin(pd.date_range(start='2005-1-6',end='2022-1-6',freq=pd.DateOffset(years=1)))]
```

|            | Open       | High       | Low        | Close      | Adj Close  | Volume   | month_name | weekday_name |
|------------|------------|------------|------------|------------|------------|----------|------------|--------------|
| Date       |            |            |            |            |            |          |            |              |
| 2005-01-06 | 97.175758  | 97.584229  | 93.509506  | 93.922951  | 93.922951  | 20852067 | January    | Thursda      |
| 2006-01-06 | 227.581970 | 234.371521 | 225.773743 | 231.960556 | 231.960556 | 35646914 | January    | Frida        |
| 2009-01-06 | 165.868286 | 169.763687 | 162.585587 | 166.406265 | 166.406265 | 12898566 | January    | Tuesda       |
| 2010-01-06 | 311.761444 | 311.761444 | 302.047852 | 302.994293 | 302.994293 | 7987226  | January    | Wednesda     |
| 2011-01-06 | 304.199799 | 308.060303 | 303.885956 | 305.604523 | 305.604523 | 4131026  | January    | Thursda      |
| 2012-01-06 | 328.344299 | 328.767700 | 323.681763 | 323.796326 | 323.796326 | 5405987  | January    | Frida        |
| 2014-01-06 | 554.426880 | 557.340942 | 551.154114 | 556.573853 | 556.573853 | 3551864  | January    | Monda        |
| 2015-01-06 | 513.589966 | 514.761719 | 499.678131 | 500.585632 | 500.585632 | 2899940  | January    | Tuesda       |
| 2016-01-06 | 730.000000 | 747.179993 | 728.919983 | 743.619995 | 743.619995 | 1947000  | January    | Wednesda     |


```
# viz a single col
google['Close'].plot()
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2a760e42e0>



```
yt['Subscribers'].plot()
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2a77a8c100>

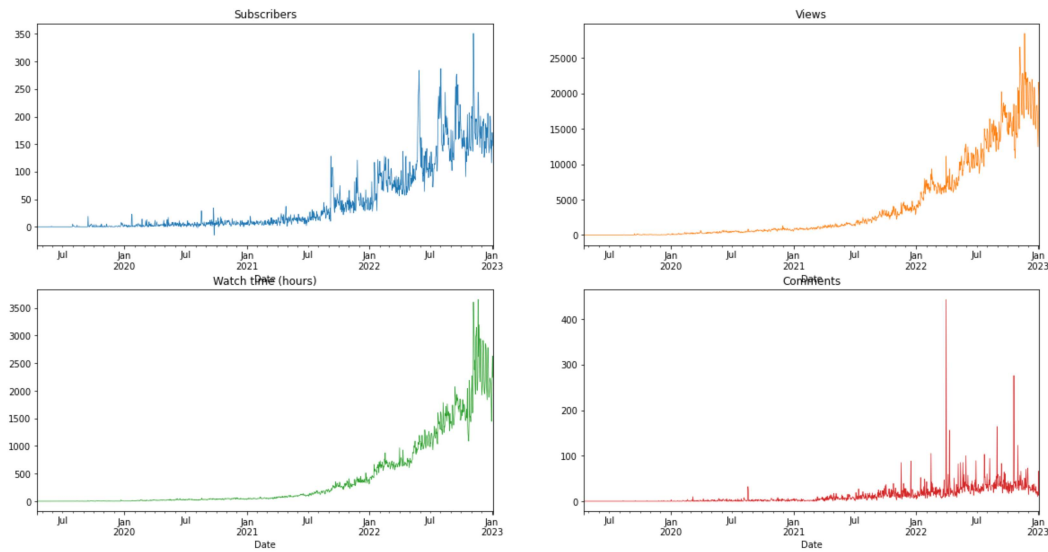


```
# viz all cols together
subset_yt = yt[['Subscribers', 'Views', 'Watch time (hours)', 'Comments added']]
ax = subset_yt.plot(subplots=True,
                    layout=(2, 2),
                    sharex=False,
```

```

sharey=False,
linewidth=0.7,
fontsize=10,
legend=False,
figsize=(20,10),
title=['Subscribers', 'Views', 'Watch time (hours)', 'Comments'])

```




```

# plot for a particular year/month/week
yt.loc['2022-12']['Subscribers'].plot()

```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2a77d25eb0>



```

google.loc['2021-12']['Close'].plot()

```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2a77d25700>



```
# monthly trend
```

```
yt.groupby('month_name')['Subscribers'].mean().plot(kind='bar')
```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2a75581730>



```
# yearly trend
```


```
yt.groupby('weekday_name')['Subscribers'].mean().plot(kind='bar')
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2a75504fa0>



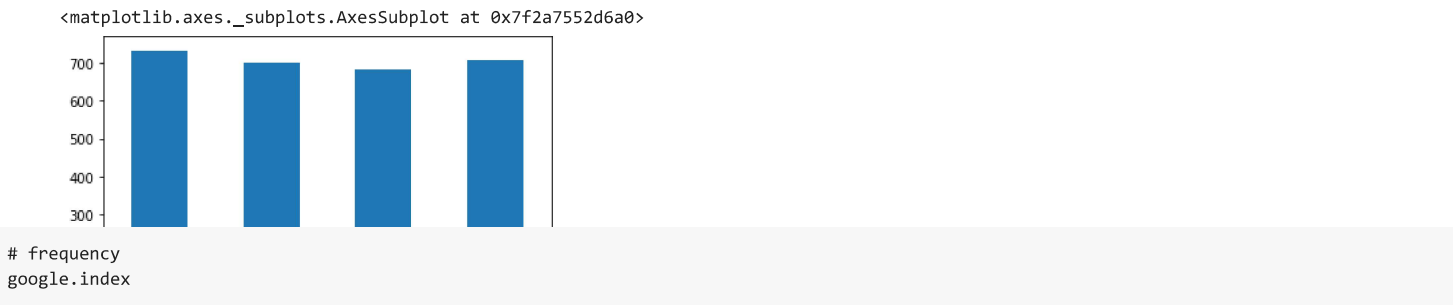
```
google.groupby('month_name')['Close'].mean().plot(kind='bar')
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2a75441a90>



```
# quarterly trend
```

```
google.groupby('quarter')['Close'].mean().plot(kind='bar')
```



```
DatetimeIndex(['2004-08-19', '2004-08-20', '2004-08-23', '2004-08-24',
               '2004-08-25', '2004-08-26', '2004-08-27', '2004-08-30',
               '2004-08-31', '2004-09-01',
               ...,
               '2022-05-09', '2022-05-10', '2022-05-11', '2022-05-12',
               '2022-05-13', '2022-05-16', '2022-05-17', '2022-05-18',
               '2022-05-19', '2022-05-20'],
              dtype='datetime64[ns]', name='Date', length=4632, freq='B')

# asfreq
google.asfreq('6H',method='bfill')
```

|                     | Open      | High      | Low       | Close     | Adj Close | Volume     | month_name | weekday |
|---------------------|-----------|-----------|-----------|-----------|-----------|------------|------------|---------|
| Date                |           |           |           |           |           |            |            |         |
| 2004-08-19 00:00:00 | 49.813290 | 51.835709 | 47.800831 | 49.982655 | 49.982655 | 44871361.0 | August     | Thu     |
| 2004-08-19 06:00:00 | 50.316402 | 54.336334 | 50.062355 | 53.952770 | 53.952770 | 22942874.0 | August     |         |
| 2004-08-19 12:00:00 | 50.316402 | 54.336334 | 50.062355 | 53.952770 | 53.952770 | 22942874.0 | August     |         |
| 2004-08-19 18:00:00 | 50.316402 | 54.336334 | 50.062355 | 53.952770 | 53.952770 | 22942874.0 | August     |         |
| 2004-08-20 00:00:00 | 50.316402 | 54.336334 | 50.062355 | 53.952770 | 53.952770 | 22942874.0 | August     |         |
| ...                 | ...       | ...       | ...       | ...       | ...       | ...        | ...        | ...     |
| 2022-05-            |           |           |           |           |           |            |            |         |

▼ Resampling

Resampling involves changing the frequency of your time series observations.

Two types of resampling are:

**Upsampling:** Where you increase the frequency of the samples, such as from minutes to seconds.

**Downsampling:** Where you decrease the frequency of the samples, such as from days to months.


```
# Downsampling
yt['Subscribers'].plot()
yt['Subscribers'].resample('Y').mean().plot()
```



```
<matplotlib.axes._subplots.AxesSubplot at 0x7f2a74e2f280>
```

```
# Upsampling
google['Close'].resample('12H').interpolate(method='spline', order=2).plot()
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f2a751e70d0>
```

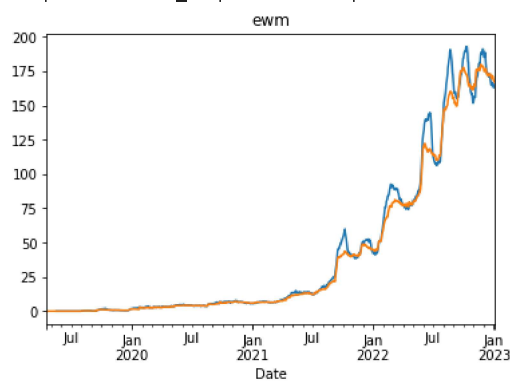


## ▼ Rolling Window(Smoothing)

Time series data in original format can be quite volatile, especially on smaller aggregation levels. The concept of rolling, or moving averages is a useful technique for smoothing time series data.

```
# Rolling window
yt['Subscribers'].rolling(30).mean().plot(title='rolling')
yt['Subscribers'].ewm(30).mean().plot(title='ewm')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f2a6d4ae9d0>
```



## ▼ Shifting

The `shift()` function in Pandas is used to, well, shift the entire series up or down by the desired number of periods.

```
# shift
yt['Subscribers']
```

```
Date
2019-04-17    0
2019-04-18    0
2019-04-19    0
2019-04-20    0
2019-04-21    0
...
2022-12-31   116
2023-01-01   142
2023-01-02   171
2023-01-03   162
2023-01-04   147
Name: Subscribers, Length: 1359, dtype: int64
```

```
yt['Subscribers'].shift(-1)
```

```
Date
2019-04-17    0.0
```

```

2019-04-18    0.0
2019-04-19    0.0
2019-04-20    0.0
2019-04-21    0.0
...
2022-12-31   142.0
2023-01-01   171.0
2023-01-02   162.0
2023-01-03   147.0
2023-01-04    NaN
Name: Subscribers, Length: 1359, dtype: float64

```

```

# shift example
df = pd.read_csv('/content/login.csv',header=None)
df = df[[1,2]]
df.head()
df.rename(columns={1:'user_id',2:'login_time'},inplace=True)
df.head()

```

|   | user_id | login_time          |
|---|---------|---------------------|
| 0 | 466     | 2017-01-07 18:24:07 |
| 1 | 466     | 2017-01-07 18:24:55 |
| 2 | 458     | 2017-01-07 18:25:18 |
| 3 | 458     | 2017-01-07 18:26:21 |
| 4 | 592     | 2017-01-07 19:09:59 |

```

user_df = df[df['user_id'] == 458]
user_df.head()

```

|    | user_id | login_time          |
|----|---------|---------------------|
| 2  | 458     | 2017-01-07 18:25:18 |
| 3  | 458     | 2017-01-07 18:26:21 |
| 9  | 458     | 2017-01-09 11:13:12 |
| 10 | 458     | 2017-01-09 11:34:02 |
| 25 | 458     | 2017-01-10 12:14:11 |

```

user_df['login_time'] = pd.to_datetime(user_df['login_time'])
user_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 208 entries, 2 to 1018
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   user_id     208 non-null    int64
1   login_time  208 non-null    datetime64[ns]
dtypes: datetime64[ns](1), int64(1)
memory usage: 4.9 KB
<ipython-input-269-fd41e73e6ce6>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

user_df['login_time'] = pd.to_datetime(user_df['login_time'])

```

```

user_df['shifted'] = user_df['login_time'].shift(1)
(user_df['login_time'] - user_df['shifted']).mean()

```

```

<ipython-input-270-091b95be4a6b>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

user_df['shifted'] = user_df['login_time'].shift(1)
Timedelta('0 days 17:29:22.053140096')

```

```
ax = df.plot(subplots=True,  
             layout=(3, 2),  
             sharex=False,  
             sharey=False,  
             linewidth=0.7,  
             fontsize=10,  
             legend=False,  
             figsize=(20,15),  
             title=['Open', 'High', 'Low', 'Close', 'Adjusted Close', 'Volume'])
```